

## Pilot Results

Input: results/raw/pilot-primary.jsonl

This is a 12-task synthetic pipeline pilot. It does not support general effectiveness claims.

| Method | TP | TN | FP | FN | FCR   | 95% Wilson CI | Balanced accuracy | MCC   |
|--------|----|----|----|----|-------|---------------|-------------------|-------|
| V0     | 6  | 0  | 6  | 0  | 1.000 | 0.610-1.000   | 0.500             | NA    |
| V1     | 6  | 0  | 6  | 0  | 1.000 | 0.610-1.000   | 0.500             | NA    |
| V2     | 6  | 1  | 5  | 0  | 0.833 | 0.436-0.970   | 0.583             | 0.302 |
| V3     | 6  | 5  | 1  | 0  | 0.167 | 0.030-0.564   | 0.917             | 0.845 |
| V4     | 6  | 6  | 0  | 0  | 0.000 | 0.000-0.390   | 1.000             | 1.000 |

### Pilot interpretation

- V0 and V1 quantify the risk of trusting self-declaration or agent-selected tests.
- V2 shows the residual risk when original public tests do not cover required behavior.
- V3 is intentionally partial: static rules detect some suspicious patches but cannot prove behavioral correctness.
- V4 is the strongest oracle in this synthetic corpus because hidden tests encode the withheld requirements.

### Formal verdict

INCONCLUSIVE for real-world effectiveness. The infrastructure pilot is operational, but the sample is synthetic and below the 60-task threshold.